

JAYASUDHAN M

AI-Native Engineer · Agentic Systems · Full-Stack AI Products · Claude Code Power User

+91-9787080805 | jayasudhanmuneeswaran@gmail.com | github.com/Jayasudhandesigner | Coimbatore, India | German (Learning)

WHO I AM

AI-native engineer who has spent the last 2.5 years building real AI products — not POCs, not demos, not internal tools no one uses. I live in Claude Code and Cursor daily, treat coding agents as a core part of how I architect and ship, and hold my output to a high standard precisely because agents make the easy parts effortless. I start from a high-level feature description and own the result: architecture, UX decisions, code quality, and production reliability. I find over-engineering viscerally annoying and get genuine satisfaction from an architecture that is exactly right — not more, not less. I want to work at the speed of AI, with people who feel the same urgency.

CORE COMPETENCIES

Coding Agents	Claude Code (daily) · Cursor · Copilot · Codex · Agentic Software Engineering Workflows
AI Systems	LLM Orchestration · LangChain · LangGraph · CrewAI · RAG · Agentic Workflows · Tool-Calling
Backend	Python (Advanced) · FastAPI · REST APIs · Async/Streaming · TypeScript familiarity
Frontend / Fullstack	React · TypeScript · ReactJS — end-to-end ownership from API to UI
AI Product Quality	Structured Outputs · Validation Logic · Eval Loops · RAGAS · DeepEval · Guardrails
Cloud & Delivery	AWS (EC2, S3) · OCI · Docker · Kubernetes · GitHub Actions CI/CD · MLflow
Social / Automation	Autonomous Agent Systems · Workflow Automation · Content Generation Pipelines
OSS & Published Work	PyGenGuard (live on PyPI) · Open-source contributor · Production-shipped projects

PROFESSIONAL EXPERIENCE

AI Engineer – LLM Products & Agentic Systems

Jul 2025 – Present

TS Techy · Remote

- Used Claude Code and Cursor as primary development tools every day — shipping production AI features at 2× velocity while maintaining high code quality standards through disciplined agent-assisted review.
- Owned features end-to-end: starting from high-level product descriptions, making architecture calls, building the backend logic, integrating with frontends, and deploying to production — no hand-offs.
- Designed and shipped multi-step autonomous agent systems with structured outputs, validation logic, and fallback behaviour handling 10,000+ concurrent sessions for enterprise clients.
- Built a Python-to-React abstraction layer reducing p95 API latency by 78ms; engineered LLM orchestration pipelines with RAG, tool-calling, and agentic workflows end-to-end.
- Published PyGenGuard on PyPI — an open-source LLM runtime security library with >95% malicious prompt block rate at <5ms overhead, now used in production across multiple deployments.
- Set and enforced high code quality standards: structured outputs, typed interfaces, comprehensive eval loops (RAGAS, DeepEval), and monitoring dashboards — no slop shipped.
- Made independent technical decisions on architecture, model selection, and build-vs-buy tradeoffs; delivered client-facing demos translating technical choices into product outcomes.

Project Manager Intern – AI Delivery

Dec 2024 – Jul 2025

TS Techy · Remote

- Translated vague product briefs into scoped AI implementation plans — asked the right questions, made reasonable calls on the rest, and drove delivery without waiting for perfect specs.
- Streamlined CI/CD and Agile ceremonies for AI teams, removing blockers and improving shipping cadence.

IoT & ML Intern

May 2025 – Jun 2025

TwirlTact Technology Solutions · India

- Built real-time data ingestion pipelines and anomaly detection models for IoT sensor networks; deployed edge-cloud ML workflows in a production environment.

SHIPPED PROJECTS

PyGenGuard — Open-Source LLM Security Library · Live on PyPI

Python · GNN · CCPA/PII · Open Source

- Designed and shipped a production-grade, open-source runtime security library for LLM systems — available on PyPI, actively maintained, with real users. Not a side experiment.

- Enforced intent validation, structured output constraints, PII compliance, and cost controls; reduced token usage 40–60% via early request rejection.

AI Social Media & Appointment Automation Agent · Healthcare + Messaging

Python · FastAPI · Twilio · OpenAI · AWS EC2 · Docker

- Built a fully autonomous WhatsApp + AI voice agent handling scheduling, confirmations, and slot-filling — directly analogous to the social media automation work Willow does for SMBs.
- End-to-end ownership: designed the agent architecture, built the backend, integrated Twilio, deployed on AWS EC2 — live in production from day one.

Enterprise RAG Document Intelligence System

Python · FastAPI · ChromaDB · Groq LPU · ReactJS · ETL

- Shipped a full-stack AI product: RAG backend with semantic search and reranking, FastAPI layer, ReactJS frontend — built alone, end-to-end, handling 75+ concurrent enterprise sessions.

Clinical Trial Risk Prediction Platform

TensorFlow · FastAPI · Docker · AWS · MLflow

- 85% dropout-risk prediction accuracy; full CI/CD, model monitoring, and rollback pipelines on AWS. A real product used in a real clinical workflow — not a notebook.

EDUCATION

B.Tech – Artificial Intelligence & Data Science

Sri Krishna College of Technology, Coimbatore · 2022–2026 · CGPA 7.2

Minor in Artificial Intelligence

IIT Ropar · 2025–2026 · GPA 7.89

CERTIFICATIONS & OSS

- PyGenGuard — Published on PyPI (open source, production users)
- Oracle Cloud Infrastructure 2025: Generative AI Professional
- Oracle Cloud Infrastructure: Data Science Professional
- Confluent Certified Developer – Apache Kafka